

Introduction to Machine Learning: Supervised Learning Final Project

Ocean Acidification Prediction Using Machine Learning Algorithms

A growing concern in climate change and environmental science is Ocean Acidification. The ocean absorbs around 30 percent of the world's carbon dioxide emissions. Carbon dioxide in the air is dissolved into ocean water where it is broken into hydrogen ions and bicarbonate. This reaction in turn decreases the pH of the ocean making it more acidic. An increase in acidity poses a threat to marine life and biological processes. Over the last couple of decades, the National Ocean and Atmospheric Administration has set up carbon dioxide and pH monitoring systems on research cruise ships and water buoy observation stations along coastlines to gain a greater understanding of this issue.

Being able to predict and forecast ocean acidity trends could have a useful impact on policy and research. In my project I will compare the performance of three different machine learning regression algorithms in using ocean chemistry data to predict pH levels. I will look at Linear Regression, Gradient Boosting Regression and Decision Tree Regression.

Import Libraries

The following libraries will be used in my project.

```
In [285.. #####  
# Importing Libraries used for the project  
#####  
  
%matplotlib inline  
import numpy as np  
import scipy as sp  
import scipy.stats as stats  
import pandas as pd  
import glob  
import os  
import matplotlib.pyplot as plt  
import seaborn as sns  
sns.set()  
import statsmodels.formula.api as smf  
import statsmodels.api as sm  
from sklearn.model_selection import train_test_split  
from sklearn.metrics import mean_squared_error, r2_score  
from sklearn.linear_model import LinearRegression  
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.datasets import make_regression
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.inspection import permutation_importance
```

Load Data

There are many ocean water quality sensors worldwide and many provide data freely accessible on the National Oceanic and Atmospheric Administration's website. In order to have more consistency in data and a manageable yet meaningful size, I chose one station, near La Parguera, Puerto Rico.

I pulled my data files from the NOAA's website under the National Centers for Environmental Information section. These files contained chemical, meteorological, physical and time sensor data from the La Parguera mooring located at 67W 18N in the Caribbean Sea near La Parguera, Puerto Rico.

The data I retrieved was in the form of csv files that were broken into approximately 12 month spans with several sensor readings per day. I choose to look at files starting from Dec 2012 (when the pH sensor was installed) until the most current, Oct 2021. The data totaled 5.2 MB in size. I stored the files in one folder and merged them into a single data frame. This dataframe had 29 columns and 22535 rows. The code below shows a summary of the column names and data types for the initial dataset.

The features I will be most interested in this analysis are:

Temperature- due to its ability to impact gas (such as CO₂) concentrations in water.

Salinity- the salt concentration can affect the alkalinity (basic-ness) of the water.

Dissolved Oxygen- impacts the chemical process. PH- measures the level of acidity,

ranges from 0-14 and lower numbers are more acidic. Partial pressure of CO₂- relates to the concentration of CO₂ in water or air. Chlorophyll- also has the ability to absorb carbon dioxide.

Data Source Citation:

Sutton, Adrienne J.; Sabine, Christopher L.; Morell, Julio M.; Musielewicz, Sylvia; Maenner Jones, Stacy; Dietrich, Colin; Bott, Randy; Osborne, John (2014). High-resolution ocean and atmosphere pCO₂ time-series measurements from mooring La_Parguera_67W_18N in the Caribbean Sea (NCEI Accession 0117354). NOAA National Centers for Environmental Information. Dataset.
https://doi.org/10.3334/cdiac/otg.tsm_la_parguera_67w_18n. Accessed Aug 3, 2024.

```
In [317... #####
# Load CSV files and merge into one data frame
#####
```

```
df = pd.concat(map(pd.read_csv, glob.glob(os.path.join('', "ocean data/La_Pa
df.head()

print("Original Data Frame Length:", len(df))
print(df.dtypes)
```

```
Original Data Frame Length: 22525
Mooring Name          object
Latitude              float64
Longitude             float64
Date                 object
Time                 object
xC02 SW (wet) (umol/mol) float64
C02 SW QF              int64
H2O SW (mmol/mol)     float64
xC02 Air (wet) (umol/mol) float64
C02 Air QF             int64
H2O Air (mmol/mol)    float64
Licor Atm Pressure (hPa) float64
Licor Temp (C)        float64
MAPC02 %O2            float64
SST (C)               float64
Salinity              float64
xC02 SW (dry) (umol/mol) float64
xC02 Air (dry) (umol/mol) float64
fC02 SW (sat) (uatm)   float64
fC02 Air (sat) (uatm)  float64
dfC02                 float64
pC02 SW (sat) (uatm)   float64
pC02 Air (sat) (uatm)  float64
dpC02                 float64
pH (total scale)       float64
pH QF                  int64
CHL (ug/l)             float64
CHL QF                  int64
NTU (NTU)              float64
NTU QF                  int64
DOXY (umol/kg)         float64
DOXY QF                 int64
dtype: object
```

Clean Data

I removed several columns such as location, date/time, mooring name, and other features that seemed redundant or unnecessary for modeling pH. Missing values were encoded as -999. I convert these to na's and then dropped any rows that contained na's. This had a drastic impact on the number of rows which brought it down to 2197. However, this is still large enough for modeling purposes. I also cleaned up column names in order to make it easier to work with. I viewed the statistical description of the data and created boxplot to verify significant outliers didn't exist. The final data frame contained 2197 rows with the following columns:

Variable	Description
----------	-------------

Variable	Description
pCO2_SW	Partial Pressure of CO2 in sea water
pCO2_Air	CO2 of air four feet above sea level
SST	Sea Surface Temperature
Salinity	Portion of salt in water
CHL	Total Chlorophyll
DOXY	Dissolved Oxygen
pH	pH measurement of ocean water

```
In [287... #####
# Cleaning the data
#####

# replace bad value with na
df = df.replace(to_replace = -999, value = np.nan)

# removed date, location, mooring, etc
# selecting only features needed for project
# and renaming columns

df_new = df[['pCO2 SW (sat) (uatm)', 'pCO2 Air (sat) (uatm)', 'SST (C)', 'Salinity (psu)'])
df_new.rename(columns={'pCO2 SW (sat) (uatm)': 'pCO2_SW', 'pCO2 Air (sat) (uatm)': 'pCO2_Air', 'SST (C)': 'SST', 'Salinity (psu)': 'Salinity'})

#drop any na's
df_new = df_new.dropna()
print("cleaned length", len(df_new))

#print stats summary to catch any usual data points
print(df_new.describe())

# look for outliers with boxplot
fig, axes = plt.subplots(2, 3, figsize=(10, 5))

# Plot on the first subplot
sns.boxplot(data=df_new, x="pCO2_SW", ax=axes[0,0])
sns.boxplot(data=df_new, x="pCO2_Air", ax=axes[0,1])
sns.boxplot(data=df_new, x="SST", ax=axes[0,2])
sns.boxplot(data=df_new, x="Salinity", ax=axes[1,0])
sns.boxplot(data=df_new, x="CHL", ax=axes[1,1])
sns.boxplot(data=df_new, x="DOXY", ax=axes[1,2])

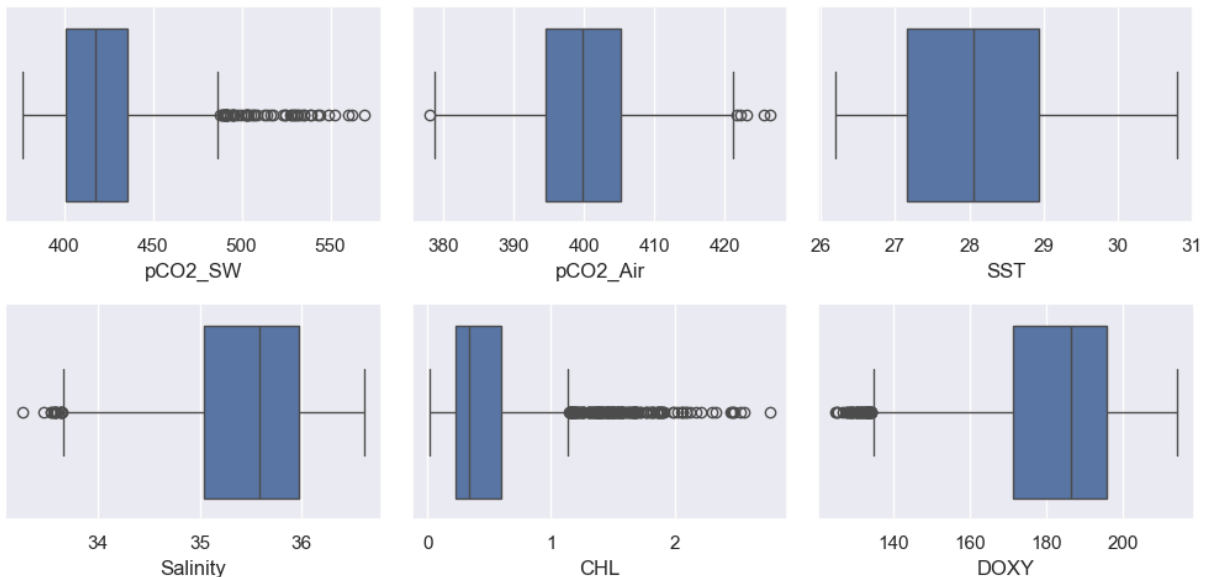
plt.tight_layout()

# show the plot
plt.show()
```

cleaned length 1569

	pCO2_SW	pCO2_Air	SST	Salinity	pH
count	1569.000000	1569.000000	1569.000000	1569.000000	1569.000000
mean	423.026769	399.988082	28.071955	35.464212	8.024362
std	29.737656	7.702659	1.056878	0.690416	0.029680
min	376.500000	377.900000	26.196000	33.261000	7.889000
25%	401.000000	394.500000	27.159000	35.044000	8.005000
50%	417.800000	399.800000	28.056000	35.589000	8.030000
75%	435.600000	405.200000	28.944000	35.971000	8.048000
max	569.000000	426.600000	30.803000	36.613000	8.092000

	CHL	DOXY
count	1569.000000	1569.000000
mean	0.501581	180.855131
std	0.442574	19.956333
min	0.012000	125.000000
25%	0.225000	171.400000
50%	0.332000	186.400000
75%	0.594000	195.800000
max	2.778000	214.300000



Exploratory Data Analysis


In my exploratory data analysis, I created correlation graphs and pairwise bivariate plots. The correlation graph shows the strongest correlation being between the pCO2_SW and pH (-0.84). Another strong correlation is pCO2_SW and SST (0.73) so this will be taken into consideration in the model. DOXY also shows some correlation with pH (0.53). The pairwise plot also shows a strong negative relationship between the partial pressure of the sea water (pCO2_SW) and the pH. As seawater CO2 increases, pH (acidity) decreases. A similar trend is seen with surface temperature (SST). The dissolved oxygen has a positive relationship. As dissolved O2 increased, pH increases. It's a little hard to see significant relationships between other features.

In [189... #####
EDA

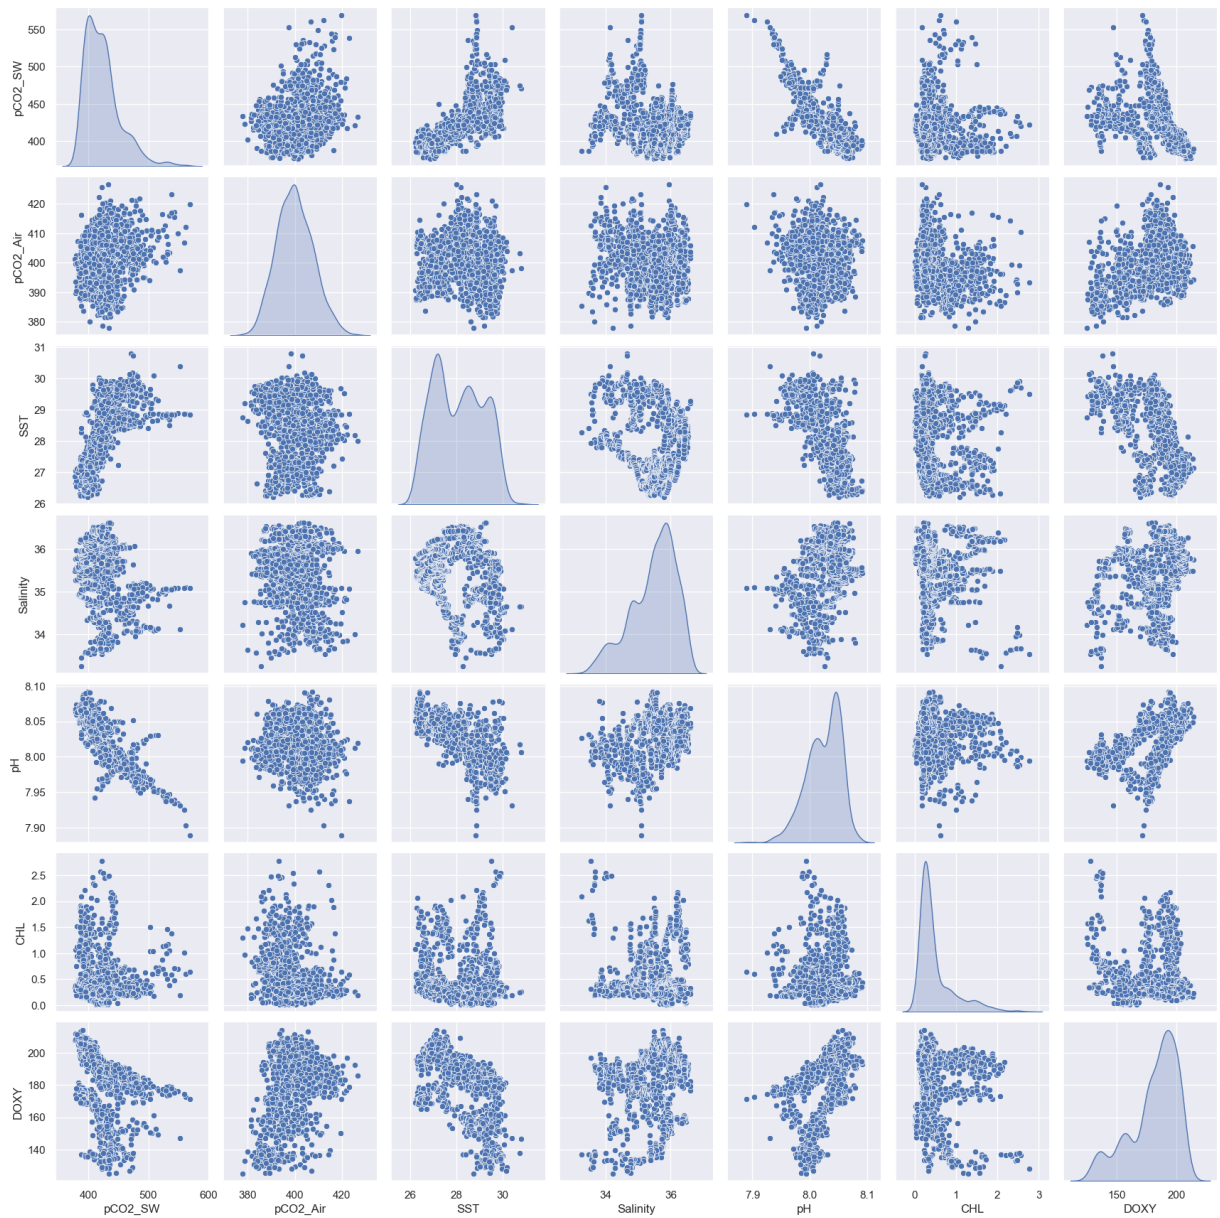
```
#####

# Correlation Matrix
matrix = df_new.corr().round(2)
plt.figure(figsize = (10,10))
heatmap = sns.heatmap(matrix, annot=True, cmap = 'vlag')
heatmap.xaxis.tick_top()

# Pairwise plot
plt.figure(figsize = (20,25))
pairplot = sns.pairplot(df_new, diag_kind = 'kde')
```



<Figure size 2000x2500 with 0 Axes>



Models

In this project I will compare three different types of regression models to make predictions. Before running the models I did an 80/20 split of the data into train and test datasets. Below I create regression models using statsmodels and sklearn. I use Linear, Decision Tree and Gradient Boost Regression algorithms.

Linear Regression

I initially ran a linear regression model for all features. In the full model, the CHL feature had a p-value > 0.05 and can be considered insignificant. I removed CHL and also the SST feature due to the collinearity seen in the EDA above. These were removed from the test and training sets and were not used for the proceeding algorithms as well (done in next block of code). I looked at several model options with the remaining features and compared the adjusted R squared values. The model with the best adjusted R squared

was one that included the features: pCO2_SW, pCO2_Air, Salinity and DOXY. I then plotted partial regression plots that model relationships between pH and a single parameter while keeping the parameters constant.

```
In [340... #####
# Modeling
#####

#create test and training data
train_ocean, test_ocean = train_test_split(df_new, train_size = 0.8, random_

#split y (pH) data apart
train_ocean_y = train_ocean["pH"]
train_ocean_x = train_ocean.drop("pH", axis=1)

test_ocean_y = test_ocean["pH"]
test_ocean_x = test_ocean.drop("pH", axis=1)

#####
# Linear Regression
#####

model_full = smf.ols(formula = 'pH ~ pCO2_SW+pCO2_Air+SST+Salinity+CHL+DOXY')
print(model_full.summary())
print()

# Calculate Adjusted R-squared for different model options
print("Full Model R2:" ,model_full.rsquared_adj)

model_1 = smf.ols(formula = 'pH ~ pCO2_SW+pCO2_Air+Salinity+DOXY', data=train_ocean)
print("Model 1 Adj R2: ", model_1.rsquared_adj)
model_2 = smf.ols(formula = 'pH ~ pCO2_SW+Salinity+DOXY', data=train_ocean).fit()
print("Model 2 Adj R2: ", model_2.rsquared_adj)
model_3 = smf.ols(formula = 'pH ~ pCO2_SW+Salinity', data=train_ocean).fit()
print("Model 3 Adj R2: ", model_3.rsquared_adj)
model_4 = smf.ols(formula = 'pH ~ pCO2_SW', data=train_ocean).fit()
print("Model 4 Adj R2: ", model_4.rsquared_adj)
model_5 = smf.ols(formula = 'pH ~ pCO2_SW+DOXY', data=train_ocean).fit()
print("Model 5 Adj R2: ", model_5.rsquared_adj)
model_6 = smf.ols(formula = 'pH ~ pCO2_SW+pCO2_Air+DOXY', data=train_ocean).fit()
print("Model 6 Adj R2: ", model_6.rsquared_adj)

x = ["Full Model", "Model 1", "Model 2", "Model 3", "Model 4", "Model 5", "Model 6"]
adjr2_train = [model_full.rsquared_adj, model_1.rsquared_adj, model_2.rsquared_adj, model_3.rsquared_adj, model_4.rsquared_adj, model_5.rsquared_adj, model_6.rsquared_adj]

# Plotting Different Model Adj R squared
plt.plot(x, adjr2_train)
plt.ylabel("Adjusted R^2")
plt.title("Adjusted R^2 vs Model")
plt.show()
```



```

# Using Best Model to Predict and Evaluate Test Data
y_lm_pred_train = model_1.predict(train_ocean_x)
y_lm_pred_test = model_1.predict(test_ocean_x)

# Calculate R-squared
r2_lm_train = r2_score(train_ocean_y , y_lm_pred_train)
r2_lm_test = r2_score(test_ocean_y , y_lm_pred_test)

# Calculate Mean Squared Error
mse_lm_train = mean_squared_error(train_ocean_y, y_lm_pred_train)
mse_lm_test = mean_squared_error(test_ocean_y, y_lm_pred_test)

#Print Chosen model summary
print(model_1.summary())

#Plotting Partial Regression Plots
fig = sm.graphics.plot_partregress_grid(model_1)
fig.tight_layout(pad=1.0)

print()
print()
print("Linear Regression:")
print('Train Mean Squared Error:', mse_lm_train)
print('Train R-squared:', r2_lm_train)
print('Test Mean Squared Error:', mse_lm_test)
print('Test R-squared:', r2_lm_test)

```

OLS Regression Results

```
=====
=
Dep. Variable:          pH    R-squared:          0.77
7
Model:                  OLS    Adj. R-squared:       0.77
6
Method:                 Least Squares    F-statistic:       725.
1
Date:                   Thu, 15 Aug 2024    Prob (F-statistic): 0.0
0
Time:                   08:50:24    Log-Likelihood:    3584.
0
No. Observations:      1255    AIC:               -715
4.
Df Residuals:          1248    BIC:               -711
8.
Df Model:               6
```

Covariance Type: nonrobust

```
=====
=
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
-
Intercept      8.0991      0.036     223.136     0.000      8.028      8.17
0
pCO2_SW        -0.0007     2.19e-05    -30.048     0.000     -0.001     -0.00
1
pCO2_Air        0.0002     6.4e-05     2.858     0.004     5.73e-05     0.00
0
SST            -0.0039      0.001     -5.973     0.000     -0.005     -0.00
3
Salinity        0.0058      0.001      9.320     0.000      0.005      0.00
7
CHL            -0.0016      0.001     -1.701     0.089     -0.003      0.00
0
DOXY            0.0002     2.94e-05     6.917     0.000      0.000      0.00
0
```

```
=====
=
Omnibus:          295.455    Durbin-Watson:       2.02
2
Prob(Omnibus):    0.000    Jarque-Bera (JB):    2825.88
5
Skew:             0.804    Prob(JB):            0.0
0
Kurtosis:         10.173    Cond. No.            5.63e+0
4
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.63e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Full Model R2: 0.7760052190380093

Model 1 Adj R2: 0.76972923556961

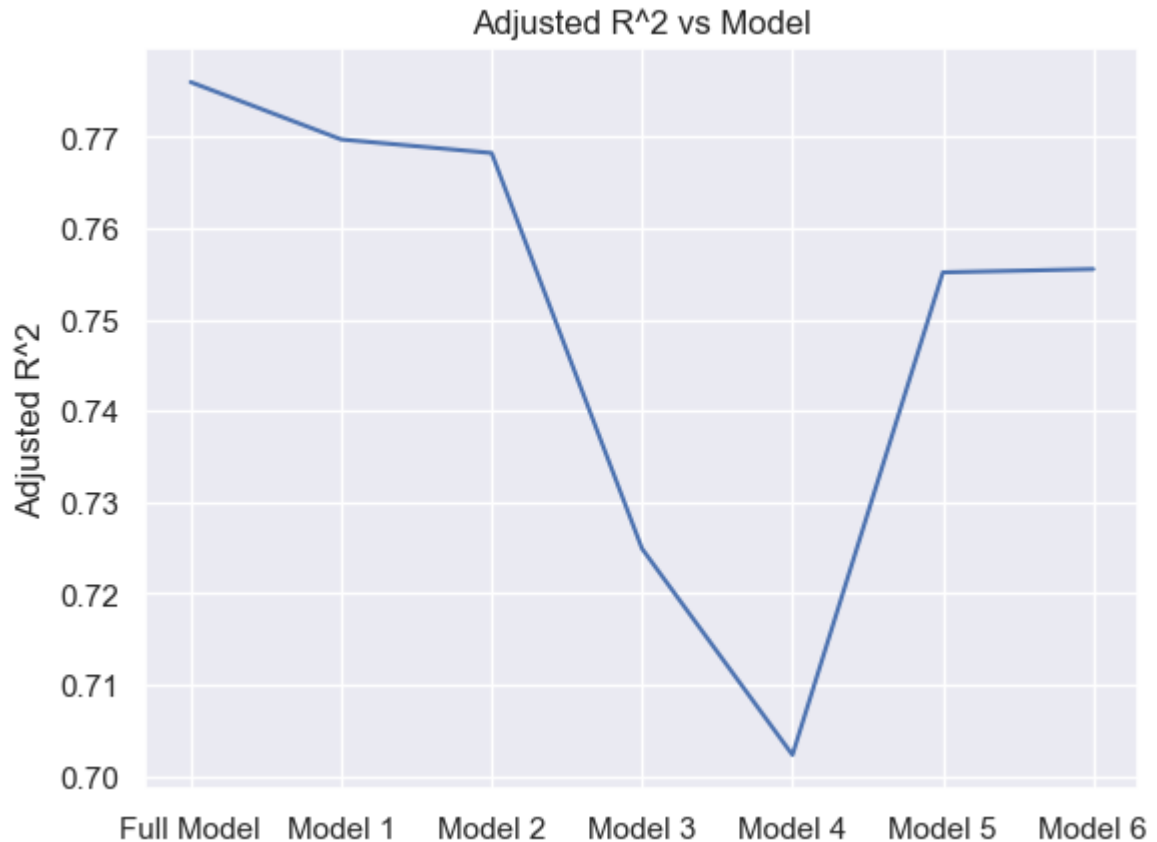
Model 2 Adj R2: 0.7682432708818645

Model 3 Adj R2: 0.7249098774922522

Model 4 Adj R2: 0.7023277392157959

Model 5 Adj R2: 0.7551708150835477

Model 6 Adj R2: 0.7555312941959244



OLS Regression Results

```

=====
=
Dep. Variable:          pH    R-squared:                0.77
0
Model:                  OLS    Adj. R-squared:           0.77
0
Method:                 Least Squares    F-statistic:           104
9.
Date:                   Thu, 15 Aug 2024    Prob (F-statistic):      0.0
0
Time:                   08:50:24    Log-Likelihood:          3565.
6
No. Observations:       1255    AIC:                     -712
1.
Df Residuals:           1250    BIC:                     -709
6.
Df Model:                4
Covariance Type:        nonrobust

```

```

=====
=
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
-
Intercept      8.0114      0.033     239.581     0.000      7.946      8.07
7
pCO2_SW        -0.0007     1.69e-05    -43.665     0.000     -0.001     -0.00
1
pCO2_Air        0.0002     6.39e-05     3.012     0.003     6.71e-05     0.00
0
Salinity        0.0055      0.001      8.839     0.000      0.004      0.00
7
DOXY            0.0003     2.59e-05    11.329     0.000      0.000      0.00
0

```

```

=====
=
Omnibus:           263.804    Durbin-Watson:           2.03
0
Prob(Omnibus):     0.000    Jarque-Bera (JB):        3210.64
0
Skew:              0.603    Prob(JB):                0.0
0
Kurtosis:          10.742    Cond. No.                5.11e+0
4

```

```

=====
=
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correc
tly specified.
[2] The condition number is large, 5.11e+04. This might indicate that there a
re

```

strong multicollinearity or other numerical problems.

Linear Regression:

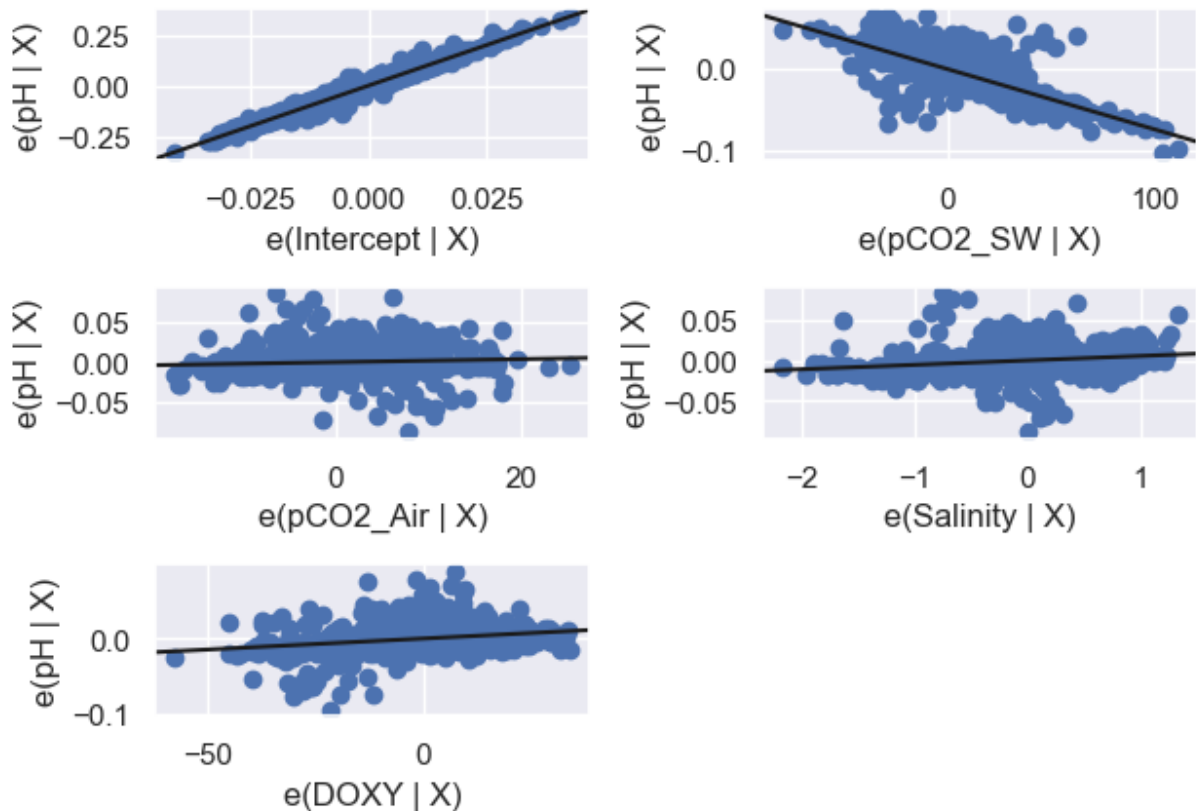
Train Mean Squared Error: 0.00019940616413150898

Train R-squared: 0.7704637515646034

Test Mean Squared Error: 0.00020546814728600324

Test R-squared: 0.7766998490460703

Partial Regression Plot



Decision Tree Regression

Next I created a Decision Tree Regression model. I set a max depth of 5 and used the default criterion of squared error. I calculated the mean squared error for the training and testing datasets. I also created Feature and Permutation Importance charts.

```
In [314... #####
# Decision Tree Regression
#####

## drop SST and CHL due to colinearity and low significance
##
train_ocean_x = train_ocean_x.drop(['SST', 'CHL'], axis=1)
test_ocean_x = test_ocean_x.drop(['SST', 'CHL'], axis=1)

dtr = tree.DecisionTreeRegressor(max_depth=5)
dtr = dtr.fit(train_ocean_x, train_ocean_y)
```

```

# Predict using train/test data
y_dtr_pred_train = dtr.predict(train_ocean_x)
y_dtr_pred_test = dtr.predict(test_ocean_x)

# Calculate Mean Squared Error
mse_dtr_train = mean_squared_error(train_ocean_y , y_dtr_pred_train)
mse_dtr_test = mean_squared_error(test_ocean_y, y_dtr_pred_test)

# Calculate R-squared
r2_dtr_train = r2_score(train_ocean_y , y_dtr_pred_train)
r2_dtr_test = r2_score(test_ocean_y , y_dtr_pred_test)

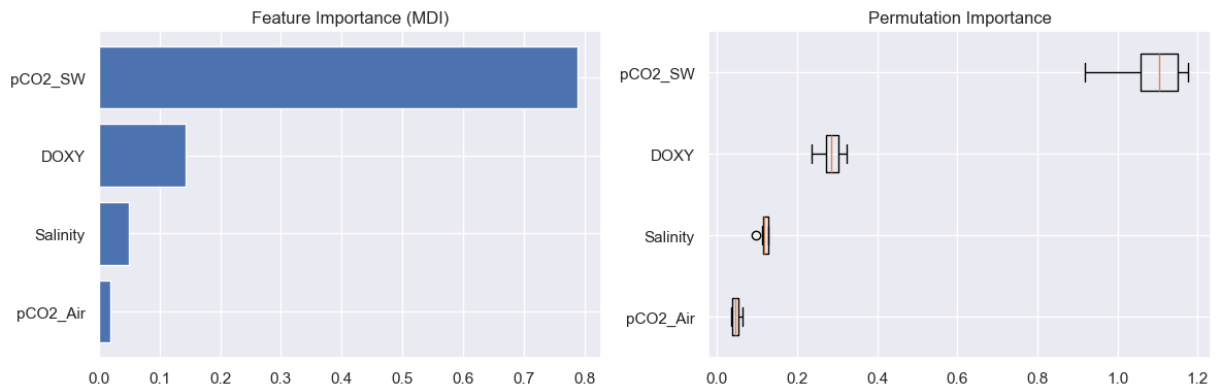
# Graph Feature and Permutation Importance
feature_importance_dtr = dtr.feature_importances_
sorted_idx_dtr = np.argsort(feature_importance_dtr)
pos_dtr = np.arange(sorted_idx_dtr.shape[0]) + 0.5
fig = plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)

plt.barh(pos_dtr, feature_importance_dtr[sorted_idx_dtr], align="center")
plt.yticks(pos_dtr, np.array(train_ocean_x.columns)[sorted_idx_dtr])
plt.title("Feature Importance (MDI)")

result_dtr = permutation_importance(
    dtr, test_ocean_x, test_ocean_y, n_repeats=10, random_state=42, n_jobs=2
)
sorted_idx_dtr = result_dtr.importances_mean.argsort()
plt.subplot(1, 2, 2)
plt.boxplot(
    result.importances[sorted_idx_dtr].T,
    vert=False,
    tick_labels=np.array(train_ocean_x.columns)[sorted_idx_dtr],
)
plt.title("Permutation Importance")
fig.tight_layout()
plt.show()

print("Decision Tree Regression:")
print('Train Mean Squared Error:', mse_dtr_train)
print('Train R-squared:', r2_dtr_train)
print('Test Mean Squared Error:', mse_dtr_test)
print('Test R-squared:', r2_dtr_test)

```



Decision Tree Regression:

Train Mean Squared Error: 9.398634562506831e-05

Train R-squared: 0.8918124057353458

Test Mean Squared Error: 0.00023353813142905152

Test R-squared: 0.7461937497834324

Gradient Boost Regression

Finally, I used a similar approach with a Gradient Boost Regression model. I computed mean squared error and r squared values for the model and created Feature and Permutation Importance charts.

```
In [315... #####
# Gradient Boost Regressor
#####

reg = GradientBoostingRegressor(random_state=0)
reg = reg.fit(train_ocean_x, train_ocean_y)

y_gbr_pred_test = reg.predict(test_ocean_x)
y_gbr_pred_train = reg.predict(train_ocean_x)

#Calculate Mean Squared Error
mse_gbr_train = mean_squared_error(train_ocean_y, y_gbr_pred_train)
mse_gbr_test = mean_squared_error(test_ocean_y, y_gbr_pred_test)

# Calculate R-squared
r2_gbr_train = r2_score(train_ocean_y, y_gbr_pred_train)
r2_gbr_test = r2_score(test_ocean_y, y_gbr_pred_test)

print("Gradient Boost Regressor: ")
print('Train Mean Squared Error:', mse_gbr_train)
print('Train R-squared:', r2_gbr_train)
print('Test Mean Squared Error:', mse_gbr_test)
print('Test R-squared:', r2_gbr_test)

feature_importance = reg.feature_importances_

sorted_idx = np.argsort(feature_importance)
pos = np.arange(sorted_idx.shape[0]) + 0.5
fig = plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
```

```
plt.barh(pos, feature_importance[sorted_idx], align="center")
plt.yticks(pos, np.array(train_ocean_x.columns)[sorted_idx])
plt.title("Feature Importance (MDI)")

result = permutation_importance(
    reg, test_ocean_x, test_ocean_y, n_repeats=10, random_state=42, n_jobs=2
)
sorted_idx = result.importances_mean.argsort()
plt.subplot(1, 2, 2)
plt.boxplot(
    result.importances[sorted_idx].T,
    vert=False,
    tick_labels=np.array(train_ocean_x.columns)[sorted_idx],
)
plt.title("Permutation Importance")
fig.tight_layout()
plt.show()
```

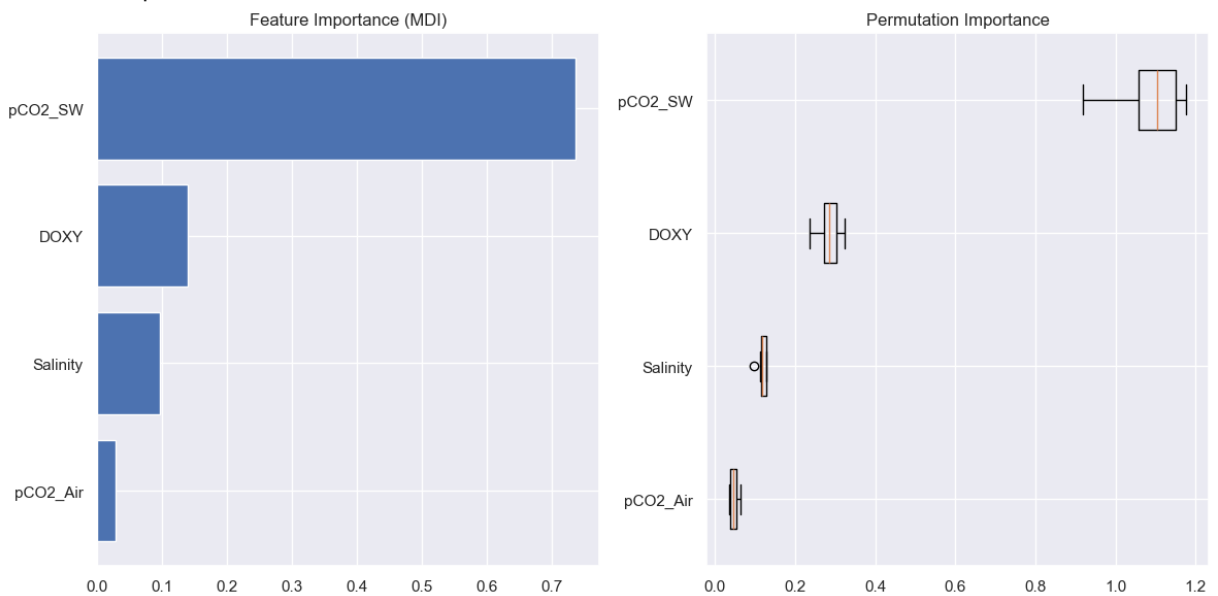
Gradient Boost Regressor:

Train Mean Squared Error: 5.7963666148998377e-05

Train R-squared: 0.9332780782813297

Test Mean Squared Error: 0.00012132408609850814

Test R-squared: 0.8681465370764554



Results and Analysis

Looking at the Partial Regression plots from the linear model, it is clear that the feature that has the greatest impact on pH in the model is the sea water partial pressure (pCO2_SW). There is a greater slope in the regression line as CO2_SW increases. The Decision Tree and Gradient Boost Regression plots for feature and permutation importance also both indicate the strength of the relationship in the model. Both show dissolved oxygen (DOXY) to be the next greatest feature of importance.

The metrics used in comparing the three algorithms were mean squared error and R squared score. These are good metrics to use when evaluating regression model performance. Mean squared values were computed for test and train sets for each

algorithm. All algorithms performed decently well, however, the Gradient Boost Regression algorithm did the best, it had the lowest mean squared error and highest R squared values. The code below graphs this information.

Mean Squared Error

	Train	Test
Linear Regression	0.00019	0.00021
Decision Tree Regression	0.00009	0.00023
Gradient Boost Regression	0.00006	0.00012

R-Squared

	Train	Test
Linear Regression	0.7704	0.7767
Decision Tree Regression	0.8918	0.7462
Gradient Boost Regression	0.9333	0.8681

Due to the observation that the sea water CO2 partial pressure showed such importance in the models, I went back and created a single parameter model using the three algorithms. Below are the performance metrics:

Mean Squared Error

	Train	Test
Linear Regression	0.00026	0.00026
Decision Tree Regression	0.00021	0.00031
Gradient Boost Regression	0.00019	0.00030

R-Squared

	Train	Test
Linear Regression	0.7026	0.7179
Decision Tree Regression	0.7551	0.6634
Gradient Boost Regression	0.7767	0.6762

The performance was not as good with the single parameter model, and it is a little less obvious which algorithm had the optimum performance. Looking at the test data, the Linear Regression algorithm gave the lowest MSE and the highest R-squared value. The code and a comparison chart for this analysis are below.

```
In [334... #####  
# Create Bar Chart for MSE and R squared Comparison  
#####
```

```

# Store Metrics from 3 models for test and train
algo = ["Linear Regression", "Decision Tree Regression", "Gradient Boost Regre
mse_test = [mse_lm_test, mse_dtr_test, mse_gbr_test]
mse_train = [mse_lm_train, mse_dtr_train, mse_gbr_train]
r2_test = [r2_lm_test, r2_dtr_test, r2_gbr_test]
r2_train = [r2_lm_train, r2_dtr_train, r2_gbr_train]

# Create Bar Chart for MSE
# set width of bar
barWidth = 0.25
fig = plt.subplots(figsize =(8, 4))

# Set position of bar on X axis
br1 = np.arange(len(mse_test))
br2 = [x + barWidth for x in br1]

# Make the plot
plt.bar(br1, mse_train, color ='r', width = barWidth,
        edgecolor ='grey', label ='MSE Train')
plt.bar(br2, mse_test, color ='g', width = barWidth,
        edgecolor ='grey', label ='MSE Test')

# Adding Labels
plt.ylabel('Mean Squared Error', fontweight ='bold', fontsize = 15)
plt.xticks([r + barWidth for r in range(len(mse_test))], algo)
plt.title("Mean Square Error for Three Different Machine Learning Algorithms")

plt.legend()
plt.show()

# Create Bar Chart for R squared

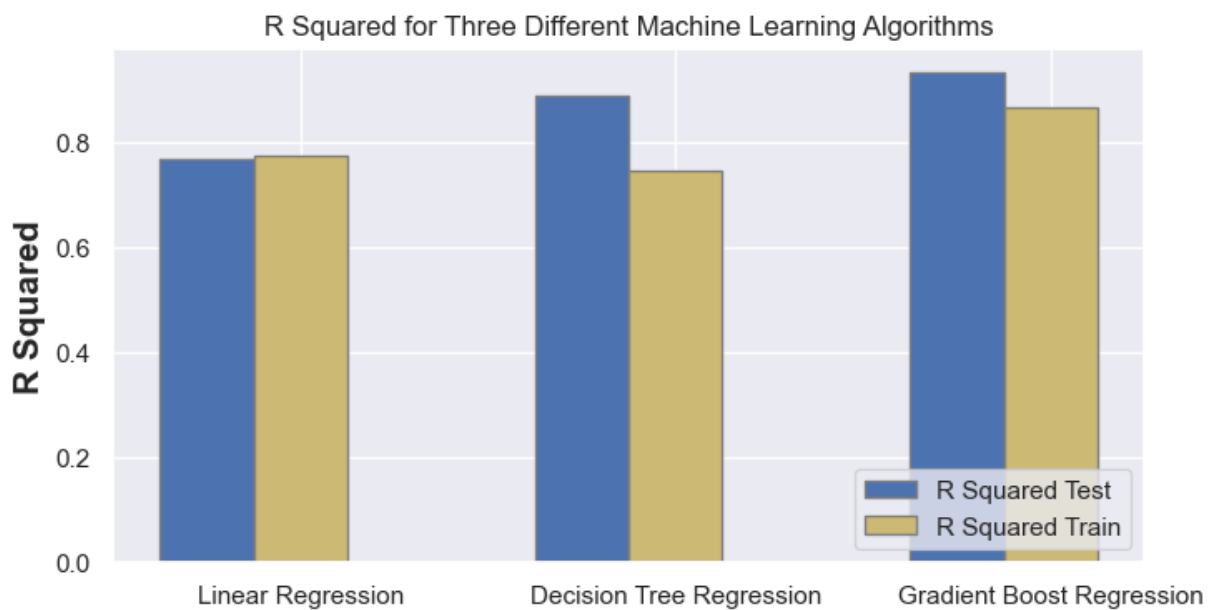
# set width of bar
barWidth = 0.25
fig = plt.subplots(figsize =(8, 4))

# Set position of bar on X axis
br1 = np.arange(len(r2_test))
br2 = [x + barWidth for x in br1]

# Make the plot
plt.bar(br1, r2_train, color ='b', width = barWidth,
        edgecolor ='grey', label ='R Squared Train')
plt.bar(br2, r2_test, color ='y', width = barWidth,
        edgecolor ='grey', label ='R Squared Test')

# Adding Labels
plt.ylabel('R Squared', fontweight ='bold', fontsize = 15)
plt.xticks([r + barWidth for r in range(len(r2_test))], algo)
plt.title("R Squared for Three Different Machine Learning Algorithms")
plt.legend(loc='lower right')
plt.show()

```



```
In [349... #####
# Running models with only pCO2_SW parameter
#####

train_ocean_x_pCO2SW = train_ocean_x['pCO2_SW']
test_ocean_x_pCO2SW = test_ocean_x['pCO2_SW']

#plt.scatter(train_ocean_x_pCO2SW, train_ocean_y)
# Linear Reg
lm_pCO2SW = smf.ols(formula = 'pH ~ pCO2_SW', data=train_ocean).fit()
print(lm_pCO2SW.summary())

# LM Predict
y_lm_pred_train_pCO2SW = lm_pCO2SW.predict(train_ocean_x_pCO2SW)
y_lm_pred_test_pCO2SW = lm_pCO2SW.predict(test_ocean_x_pCO2SW)

train_ocean_x_pCO2SW = train_ocean_x_pCO2SW.to_numpy()
test_ocean_x_pCO2SW = test_ocean_x_pCO2SW.to_numpy()
```

```

# LM R-squared / Mean Squared Error
r2_lm_train_pCO2SW = r2_score(train_ocean_y , y_lm_pred_train_pCO2SW)
r2_lm_test_pCO2SW = r2_score(test_ocean_y , y_lm_pred_test_pCO2SW)
mse_lm_train_pCO2SW = mean_squared_error(train_ocean_y, y_lm_pred_train_pCO2SW)
mse_lm_test_pCO2SW = mean_squared_error(test_ocean_y, y_lm_pred_test_pCO2SW)

#Decision Tree Reg
dtr_pCO2SW = tree.DecisionTreeRegressor()
dtr_pCO2SW = dtr.fit(train_ocean_x_pCO2SW.reshape(-1, 1), train_ocean_y)

# DTR Predict
y_dtr_pred_train_pCO2SW = dtr_pCO2SW.predict(train_ocean_x_pCO2SW.reshape(-1, 1))
y_dtr_pred_test_pCO2SW = dtr_pCO2SW.predict(test_ocean_x_pCO2SW.reshape(-1, 1))

# DTR R-squared / Mean Squared Error
mse_dtr_train_pCO2SW = mean_squared_error(train_ocean_y , y_dtr_pred_train_pCO2SW)
mse_dtr_test_pCO2SW = mean_squared_error(test_ocean_y, y_dtr_pred_test_pCO2SW)
r2_dtr_train_pCO2SW = r2_score(train_ocean_y , y_dtr_pred_train_pCO2SW)
r2_dtr_test_pCO2SW = r2_score(test_ocean_y , y_dtr_pred_test_pCO2SW)

# Gradient Boost Reg
reg_pCO2SW = GradientBoostingRegressor(random_state=0)
reg_pCO2SW = reg_pCO2SW.fit(train_ocean_x_pCO2SW.reshape(-1, 1), train_ocean_y)

# GBR Predict
y_gbr_pred_train_pCO2SW = reg_pCO2SW.predict(train_ocean_x_pCO2SW.reshape(-1, 1))
y_gbr_pred_test_pCO2SW = reg_pCO2SW.predict(test_ocean_x_pCO2SW.reshape(-1, 1))

# GBE R-squared / Mean Squared Error
mse_gbr_train_pCO2SW = mean_squared_error(train_ocean_y , y_gbr_pred_train_pCO2SW)
mse_gbr_test_pCO2SW = mean_squared_error(test_ocean_y, y_gbr_pred_test_pCO2SW)
r2_gbr_train_pCO2SW = r2_score(train_ocean_y , y_gbr_pred_train_pCO2SW)
r2_gbr_test_pCO2SW = r2_score(test_ocean_y , y_gbr_pred_test_pCO2SW)

print()
print()
print("Single Parameter Performance- pCO2_SW")
print("Linear Regression")
print("MSE train/test: ", mse_lm_train_pCO2SW, mse_lm_test_pCO2SW)
print("R squared train/test: " , r2_lm_train_pCO2SW, r2_lm_test_pCO2SW)
print("Decision Tree Regression")
print("MSE train/test: ", mse_dtr_train_pCO2SW, mse_dtr_test_pCO2SW)
print("R squared train/test: " , r2_dtr_train_pCO2SW, r2_dtr_test_pCO2SW)
print("Gradient Boost Regression")
print("MSE train/test: ", mse_gbr_train_pCO2SW, mse_gbr_test_pCO2SW)
print("R squared train/test: " , r2_gbr_train_pCO2SW, r2_gbr_test_pCO2SW)

#####
# Plotting Results for single parameter model
#####

```

```
#####
# Create Bar Chart for MSE and R squared Comparison
#####

# Store Metrics from 3 models for test and train
mse_test_pC02SW = [mse_lm_test_pC02SW, mse_dtr_test_pC02SW, mse_gbr_test_pC02SW]
mse_train_pC02SW = [mse_lm_train_pC02SW, mse_dtr_train_pC02SW, mse_gbr_train_pC02SW]
r2_test_pC02SW = [r2_lm_test_pC02SW, r2_dtr_test_pC02SW, r2_gbr_test_pC02SW]
r2_train_pC02SW = [r2_lm_train_pC02SW, r2_dtr_train_pC02SW, r2_gbr_train_pC02SW]

# Create Bar Chart for MSE
# set width of bar
barWidth = 0.25
fig = plt.subplots(figsize =(8, 4))

# Set position of bar on X axis
br1 = np.arange(len(mse_test_pC02SW))
br2 = [x + barWidth for x in br1]

# Make the plot
print()
print()
plt.bar(br1, mse_train_pC02SW, color ='b', width = barWidth,
        edgecolor ='grey', label ='MSE Train')
plt.bar(br2, mse_test_pC02SW, color ='g', width = barWidth,
        edgecolor ='grey', label ='MSE Test')

# Adding Labels
plt.ylabel('Mean Squared Error', fontweight ='bold', fontsize = 15)
plt.xticks([r + barWidth for r in range(len(mse_test))], algo)
plt.suptitle("Mean Square Error for Three Different Machine Learning Algorithms")
plt.title("Single Parameter Model using pC02_SW")
plt.legend()
plt.show()

# Create Bar Chart for R squared

# set width of bar
barWidth = 0.25
fig = plt.subplots(figsize =(8, 4))

# Set position of bar on X axis
br1 = np.arange(len(r2_test_pC02SW))
br2 = [x + barWidth for x in br1]

# Make the plot
plt.bar(br1, r2_train_pC02SW, color ='r', width = barWidth,
        edgecolor ='grey', label ='R Squared Train')
plt.bar(br2, r2_test_pC02SW, color ='y', width = barWidth,
        edgecolor ='grey', label ='R Squared Test')

# Adding Labels
plt.ylabel('R Squared', fontweight ='bold', fontsize = 15)
```

```
plt.xticks([r + barWidth for r in range(len(r2_test_pC02SW))], algo)
plt.suptitle("Mean Square Error for Three Different Machine Learning Algorithms")
plt.title("Single Parameter Model using pC02_SW")
plt.legend(loc='lower right')
plt.show()
```

OLS Regression Results

```

=====
=
Dep. Variable:          pH    R-squared:                0.70
3
Model:                  OLS    Adj. R-squared:           0.70
2
Method:                 Least Squares    F-statistic:           296
0.
Date:                   Thu, 15 Aug 2024    Prob (F-statistic):      0.0
0
Time:                   09:33:03    Log-Likelihood:         3403.
0
No. Observations:      1255    AIC:                    -680
2.
Df Residuals:          1253    BIC:                    -679
2.
Df Model:               1
Covariance Type:       nonrobust

```

```

=====
=
              coef      std err          t      P>|t|      [0.025      0.97
5]
-----
-
Intercept      8.3791      0.007    1284.148      0.000      8.366      8.39
2
pCO2_SW       -0.0008    1.54e-05    -54.403      0.000     -0.001     -0.00
1
=====

```

```

=
Omnibus:          128.117    Durbin-Watson:           2.03
1
Prob(Omnibus):    0.000    Jarque-Bera (JB):        945.16
3
Skew:             0.048    Prob(JB):                5.76e-20
6
Kurtosis:         7.250    Cond. No.                 6.08e+0
3
=====
=

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

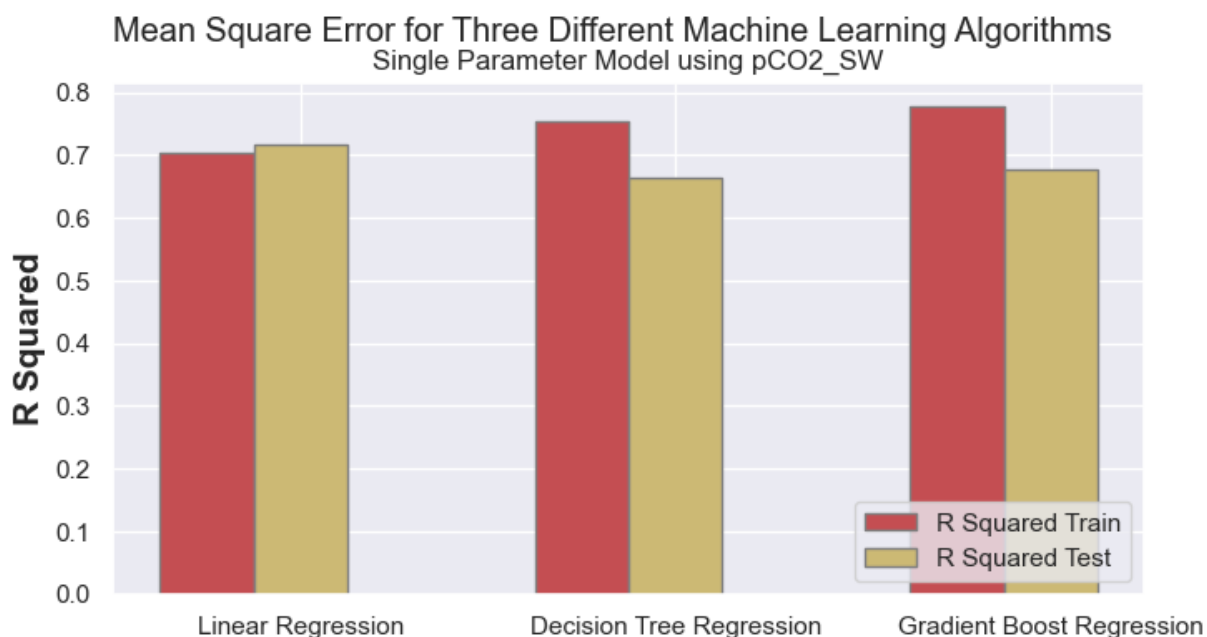
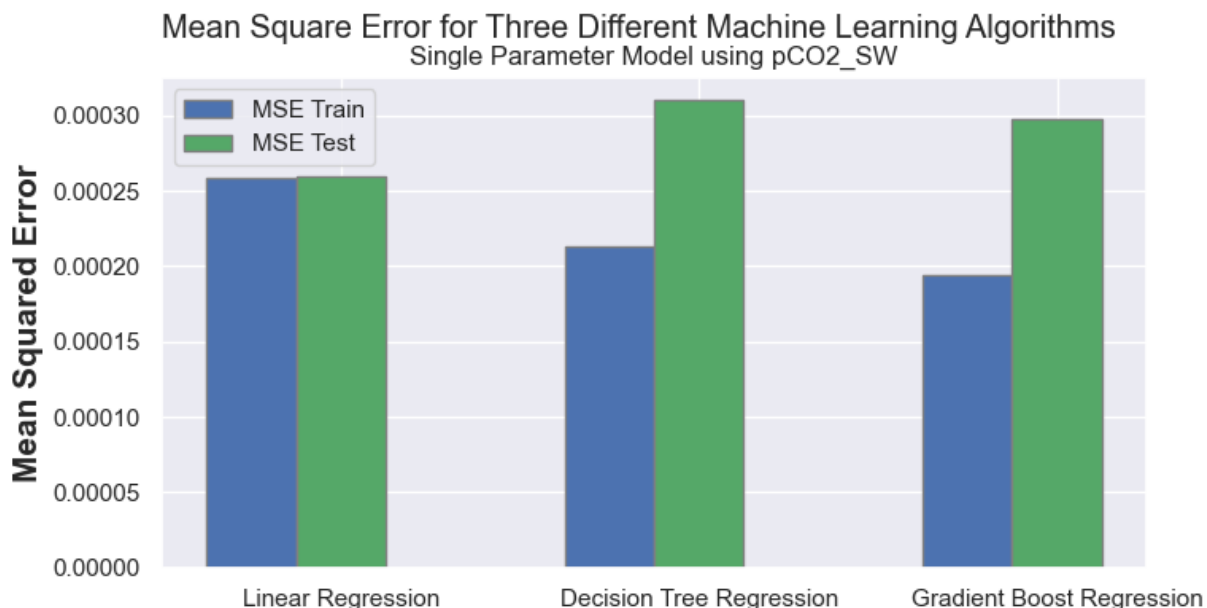
[2] The condition number is large, 6.08e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Single Parameter Performance- pCO2_SW

Linear Regression

MSE train/test: 0.00025839208151053967 0.00025953543761073555

R squared train/test: 0.7025651174141884 0.7179402103835517
 Decision Tree Regression
 MSE train/test: 0.00021272393124125016 0.0003097232527305104
 R squared train/test: 0.7551336823402125 0.6633967356877196
 Gradient Boost Regression
 MSE train/test: 0.0001940240926737965 0.00029793308304906416
 R squared train/test: 0.7766590489697519 0.6762101410958811



Conclusion

In this project I used three different machine learning algorithms to model the affect of different variables on ocean water pH. I used Linear Regression, Decision Tree Regression and Gradient Boost Regression and computed mean squared error and R squared scores as a means for evaluation. All models performed well, but the Gradient

Boost Regression algorithm showed to be very good at predicting pH levels with a multi-parametric model. The Linear Regression algorithm did well with one parameter. These algorithms could be effective tools in forecasting trends and help researchers and policy makers gain understanding in ocean acidification. In my analysis, the partial pressure of CO₂ in the sea water was shown to be the most crucial parameter for the models. For future improvements, a more scientific knowledge set of additional parameters and interaction to investigate may be needed. Another future improvement could be to find projection data for ocean water pCO₂ levels and use the models to predict changes in ocean water pH levels.

This project shows the potential and relative easiness of these algorithms to be useful tools for forecasting, whether for environmental causes or other regression prediction applications.